


Microdata Protection

Privacy Preserving Information Access: Homework 1

Reza Ghasemi

A.Y. 2022/2023



UNIVERSITÀ
DEGLI STUDI
DI PADOVA



Heart Attack Analysis & Prediction Dataset

```
In [90]: 1 heart_ds = pd.read_csv("heart.csv")
          2
          3 # display dataset
          4 heart_ds
```

Out[90]:

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1
...
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0

303 rows × 14 columns

- Used in ML to predict whether a patient is going to have heart attack or not.
- Dataset contains the following information regarding patients:

Age : Age of the patient

Sex : Sex of the patient

cp : Chest Pain type

rtbpps : resting blood pressure (in mm Hg)

chol: cholesterol in mg/dl fetched via BMI sensor

output: diagnosis of heart disease

...

- Identifiers have been removed (e.g. name, SSN). Based on the information, we see the likelihood of heart attack as *output* (medical data which must be protected).

- ♦ Fake information is generated and added

```
1 # adding names to the dataset
2 full_names = []
3
4 fake = Faker()
5
6 for _ in range(303):
7     full_names.append(fake.name())
8
9 heart_ds['names'] = full_names
10
11 # Adding marital status to the dataset
12 marital_status = []
13 possible_choices = ['single', 'married', 'divorced', 'widow']
14
15 for _ in range(303):
16     marital_status.append(random.choice(possible_choices))
17
18 heart_ds["marital status"] = marital_status
19
20 # adding race to the dataset (for simplicity I have not added other races)
21 race = []
22 existing_race = ['black', 'white', 'asian', 'native hawaiian', 'hispanic', 'american indian']
23
24 for _ in range(303):
25     race.append(random.choice(existing_race))
26
27 heart_ds["race"] = race
28
29 # adding social security number
30
31 Faker.seed(0)
32 ssn_column = []
33
34 for _ in range(303):
35     ssn_column.append(fake.ssn())
36
37 heart_ds["social security number"] = ssn_column
38
```

Expanded dataset



	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output	names	marital status	race	social security number
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1	Vicki Cummings	widow	white	865-50-6891
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1	Jay Bailey	married	asian	042-34-8377
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1	Michael Miller	single	native hawaiian	498-52-4970
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1	Christine Douglas	widow	american indian	489-46-9559
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1	Jerome Ramirez	divorced	asian	224-65-2282
...
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0	Isaac Bates	married	white	610-11-0581
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0	Alice Wilson	single	hispanic	073-34-5008
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0	James Morrow III	divorced	black	547-44-1937
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0	Tanner Cook	widow	american indian	543-32-2680
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0	Wendy Clements	married	american indian	070-54-4747

303 rows × 18 columns



Identifiers



	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output	names	marital status	race	social security number
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1		Sandra Lee	divorced	white	865-50-6891
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2		Jack Clark	divorced	asian	042-34-8377
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2		Paul Steele	married	american indian	498-52-4970
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2		Louis Ford	single	hispanic	489-46-9559
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2		Bryan Mayo	divorced	american indian	224-65-2282
...
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0	Julie Garrett	divorced	native hawaiian	610-11-0581
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0	Jonathan Nicholson	widow	american indian	073-34-5008
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0	Cynthia Mosley	divorced	white	547-44-1937
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0	Anna James	single	white	543-32-2680
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0	Robin Cooper	divorced	asian	070-54-4747

303 rows × 18 columns



Quasi-identifiers



	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output	names	marital status	race	social security number
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1	Sandra Lee	divorced	white	865-50-6891
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1	Jack Clark	divorced	asian	042-34-8377
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1	Paul Steele	married	american indian	498-52-4970
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1	Louis Ford	single	hispanic	489-46-9559
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1	Bryan Mayo	divorced	american indian	224-65-2282
...
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0	Julie Garrett	divorced	native hawaiian	610-11-0581
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0	Jonathan Nicholson	widow	american indian	073-34-5008
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0	Cynthia Mosley	divorced	white	547-44-1937
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0	Anna James	single	white	543-32-2680
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0	Robin Cooper	divorced	asian	070-54-4747

303 rows × 18 columns



Confidential Attributes



	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output	names	marital status	race	social security number
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1	Sandra Lee	divorced	white	865-50-6891
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1	Jack Clark	divorced	asian	042-34-8377
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1	Paul Steele	married	american indian	498-52-4970
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1	Louis Ford	single	hispanic	489-46-9559
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1	Bryan Mayo	divorced	american indian	224-65-2282
...
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0	Julie Garrett	divorced	native hawaiian	610-11-0581
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0	Jonathan Nicholson	widow	american indian	073-34-5008
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0	Cynthia Mosley	divorced	white	547-44-1937
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0	Anna James	single	white	543-32-2680
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0	Robin Cooper	divorced	asian	070-54-4747

303 rows × 18 columns



♦ Who releases data?

- Government agencies
- Research institutions
- Businesses (e.g. Netflix)



♦ What harms can be done?

- Reputation of business and individual
- Employability
- Personal relationships
- Physical harm



- ♦ Malicious actors
- ♦ Researchers may single out an individual by accident
- ♦ Doxxers may use information for harassment (e.g. sharing home address of target on 4chan)



The idea behind generalization is putting multiple categories under one category. (e.g. instead of divorced, widow, single and married, we have “espoused” or “not married”)


```
1 heart_ds['marital status'].unique()  
  
array(['divorced', 'married', 'single', 'widow'], dtype=object)
```

```
1 # Prepare generalization map  
2 generalization_map = {'married': 'espoused', 'widow': 'not married', 'single': 'not married', 'divorced': 'not married'}  
3  
4 heart_ds['marital_generalized'] = heart_ds['marital status'].replace(generalization_map)  
5  
6 # Display table  
7 heart_ds[['marital status', 'marital_generalized']]
```

Result:

	marital status	marital_generalized
0	divorced	not married
1	divorced	not married
2	married	espoused
3	single	not married
4	divorced	not married
...
298	divorced	not married
299	widow	not married
300	divorced	not married
301	single	not married
302	divorced	not married

303 rows × 2 columns



(n,k) dominance rule

It measures sensitivity. primary suppression should be applied when the sum of the n largest contributors exceeds $k\%$ of the cell total.

```
1 n = 3
2 k = 0.2
3 val = np.array(heart_ds[(heart_ds['race']=='black') &
4                       (heart_ds['marital status']=='divorced')]['chol'])
5
6 print(f"is the cell sensible according to {n}-{k} rule?: {nk_rule(val, n, k)}\n")
7
8 k = 0.04
9 print(f"is the cell sensible according to {n}-{k} rule?: {nk_rule(val, n, k)}\n")
10 # not sensitive because we have four individual more than 0.25
```

```
contributions: [0.06864456 0.07998837 0.06573589 0.06806283 0.0744619  0.06573589
0.0788249  0.06835369 0.0858057  0.0735893  0.05904596 0.0840605
0.06020942 0.06748109]
```

```
number of individuals contributing more than 0.2: 0
is the cell sensible according to 3-0.2 rule?: True
```

```
contributions: [0.06864456 0.07998837 0.06573589 0.06806283 0.0744619  0.06573589
0.0788249  0.06835369 0.0858057  0.0735893  0.05904596 0.0840605
0.06020942 0.06748109]
```

```
number of individuals contributing more than 0.04: 14
is the cell sensible according to 3-0.04 rule?: False
```




Domain is divided into disjointed intervals

```
In [23]: 1 heart_ds['chol_level'] = pd.cut(heart_ds['chol'], bins=10)
          2
          3 # display chol_level interval for few elements
          4 heart_ds[['chol', 'chol_level']].head()
```

Out[23]:

	chol	chol_level
0	233.0	(213.6, 257.4]
1	250.0	(213.6, 257.4]
2	204.0	(169.8, 213.6]
3	236.0	(213.6, 257.4]
4	354.0	(345.0, 388.8]



... Why not use *qcut* instead?

As we saw in the example, *qcut* may give us errors

```
1 adultdf['hpw'] = pd.qcut(adultdf['hours-per-week'], 10)
```

ValueError

Traceback (most recent call last)

Input In [18], in <cell line: 1>()

----> 1 adultdf['hpw'] = pd.qcut(adultdf['hours-per-week'], 10)

File ~/anaconda3/lib/python3.9/site-packages/pandas/core/reshape/tile.py:378, in qcut(x, q, labels, retbins, precision, duplicates)

375 x_np = x_np[~np.isnan(x_np)]

376 bins = np.quantile(x_np, quantiles)

--> 378 fac, bins = bins_to_cuts(

379 x,

380 bins,

381 labels=labels,

382 precision=precision,

383 include_lowest=True,

384 dtype=dtype,

385 duplicates=duplicates,


386)

388 return _postprocess_for_cut(fac, bins, retbins, dtype, original)

Solutions?

- ♦ **Possible Solutions:**
 - Use *cut* instead
 - Fixed since `panda >= 0.20.0` (use `duplicates='drop'`)
 - Use lower number for bins to have less number of same samples
- ♦ **Why errors?** Many samples are put in the same category while trying to maintain same quantity.
- ♦ **Issues?** outliers are not protected

Top coding and Bottom coding



Top coding:

```
1 # top coding
2 greater_index = heart_ds[heart_ds['chol'] > 240].index.to_list()
```

Bottom coding:

```
1 # bottom coding
2 below_index = heart_ds[heart_ds['chol'] < 220].index.to_list()

1 heart_ds.loc[greater_index, 'chol'] = ">240"
2 heart_ds.loc[below_index, 'chol'] = "<220"
3 heart_ds[(heart_ds['chol'] == ">240") | (heart_ds['chol'] == "<220")]
```



Top coding and Bottom coding (cont.)



We can use top coding and bottom coding to protect them

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output	names	marital status	race	social security number	marital_generalize
1	37	1	2	130	>240	0	1	187	0	3.5	0	0	2	1	Jack Clark	divorced	asian	042-34-8377	not marrie
2	41	0	1	130	<220	0	0	172	0	1.4	2	0	2	1	Paul Steele	married	american indian	498-52-4970	espouse
4	57	0	0	120	>240	0	1	163	1	0.6	2	0	2	1	Bryan Mayo	divorced	american indian	224-65-2282	not marrie
5	57	1	0	140	<220	0	1	148	0	0.4	1	0	1	1	Jocelyn Ramsey	divorced	american indian	289-18-1554	not marrie
6	56	0	1	140	>240	0	0	153	0	1.3	1	0	2	1	Kelly Rivera	married	hispanic	634-33-8726	espouse
...
297	59	1	0	164	<220	1	0	90	0	1.0	1	2	1	0	Edgar Marshall	married	native hawaiian	363-37-3212	espouse
298	57	0	0	140	>240	0	1	123	1	0.2	1	0	3	0	Julie Garrett	divorced	native hawaiian	610-11-0581	not marrie
299	45	1	3	110	>240	0	1	132	0	1.2	1	0	3	0	Jonathan Nicholson	widow	american indian	073-34-5008	not marrie
300	68	1	0	144	<220	1	1	141	0	3.4	1	2	3	0	Cynthia Mosley	divorced	white	547-44-1937	not marrie
301	57	1	0	130	<220	0	1	115	1	1.2	1	1	3	0	Anna James	single	white	543-32-2680	not marrie



A well known technique in which value of sensitive cell is removed.

Example: Removing *chol level* higher than 240 for black and divorced individuals.

```
1 # [243 198 227 265 214 172 254 174 249 231 237]
2 heart_ds.loc[(heart_ds['race']=='black')
3             & (heart_ds['marital status']=='divorced')
4             & (heart_ds['chol'] > 240), 'chol']
```

```
11      275
50      256
131     271
140     295
156     253
231     289
Name: chol, dtype: int64
```

Local Suppression (cont.)



```
1 heart_ds.loc[(heart_ds['race']=='black')
2             & (heart_ds['marital status']=='divorced')
3             & (heart_ds['chol'] > 240), 'chol'] = np.NaN
```

```
1 # Display
2
3 heart_ds.loc[(heart_ds['race']=='black')
4             & (heart_ds['marital status']=='divorced')
5             & (heart_ds['chol'] > 240), 'chol']
```

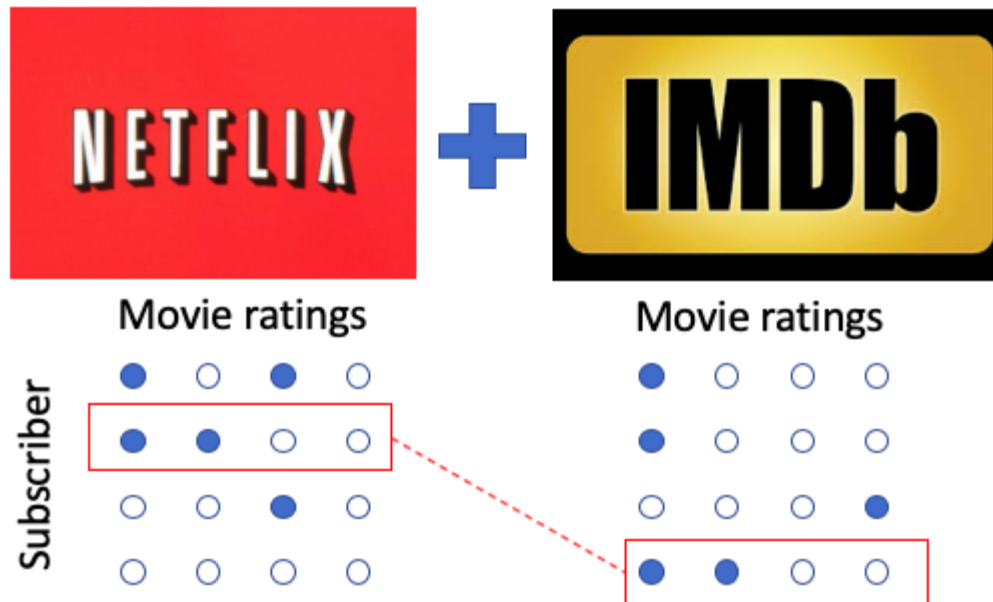
Series([], Name: chol, dtype: float64)

```
1 heart_ds.loc[(heart_ds['race']=='black')
2             & (heart_ds['marital status']=='divorced')]
```

	age	sex	cp	trtbps	chol	fbs	restecg	thalachh	exng	oldpeak	slp	caa	thall	output	names	marital status	race	social security number	marital_generalized	cl
3	56	1	1	120	236.0	0	1	178	0	0.8	2	0	2	1	Albert Montes	divorced	black	489-46-9559	not married	
11	48	0	2	130	NaN	0	1	139	0	0.2	2	0	2	1	Daniel Hale	divorced	black	363-56-5181	not married	
17	66	0	3	150	226.0	0	1	114	0	2.6	0	0	2	1	Sergio Wolfe	divorced	black	728-86-0019	not married	
49	53	0	0	138	234.0	0	0	160	0	0.0	2	0	2	1	Heather Valentine	divorced	black	120-91-3613	not married	
50	51	0	2	130	NaN	0	0	149	0	0.5	2	0	2	1	Gabriel Durham	divorced	black	381-22-5449	not married	



Two sources of information: one anonymized, another public are used for comparison (e.g. Netflix and IMDB)



Record Linkage (cont.)



```
1 # We import first dataset
2 first_dataset = pd.DataFrame(heart_ds)
3 names_from_first_list = first_d['names'].tolist() # 303 elements in list
```

```
1 # We create our second dataset (e.g. database of a bank which was leaked)
2 def generate_dataset(n):
3     output = [{"address":fake.address(),
4                 "city":fake.city(),
5                 "iban":fake.iban(),
6                 "swift":fake.swift(length=8),
7                 "cc-provider":fake.credit_card_provider()} for x in range(n)]
8     return output
9
10 second_dataset = pd.DataFrame(generate_dataset(303))
11
12 # append names from first list to second list to perform record linkage
13 second_dataset['names'] = names_from_first_list
14
15
16 # adding race to the dataset (for simplicity I have not added other races)
17 race_second = []
18 existing_race_2 = ['black', 'white', 'asian']
19
20 for _ in range(303):
21     race_second.append(random.choice(existing_race_2))
22
23 second_dataset["race"] = race_second
24
25 print(second_dataset)
```



Records are compared using *recordlinkage* module

```
1 import recordlinkage
2
3 indexer = recordlinkage.Index()
4 indexer.block("names")
5 candidate_links = indexer.index(first_dataset, second_dataset)
```

```
1 compare_cl = recordlinkage.Compare()
2
3
4 compare_cl.string("names", "names", method="jarowinkler", threshold=0.9, label="names")
5 compare_cl.string("race", "race", method="jarowinkler", threshold=0.9, label="race")
6
7 # Compare the records
8 features = compare_cl.compute(candidate_links, first_dataset, second_dataset)
9
10 # display features
11 features.describe()
```

```
1 features.sum(axis=1).value_counts().sort_index(ascending=False)
2
3 features[features.sum(axis=1) >= 0]
```

Record Linkage (cont.)



We used same name for both sets.
name is the same (1.0) while race
for some indexes is different.



*In your opinion, what are possible solutions
to prevent record linkage?*

		names	race
0	0	1.0	1.0
1	1	1.0	0.0
2	2	1.0	0.0
3	3	1.0	0.0
4	4	1.0	0.0
...
298	298	1.0	0.0
299	299	1.0	1.0
300	300	1.0	0.0
301	301	1.0	0.0
302	302	1.0	1.0

305 rows × 2 columns



- ♦ Privacy is slowly gaining attention (not just in academia)
- ♦ New libraries are beginning to be developed (e.g. PyMASq)
- ♦ Many possible attack scenarios and defenses are yet to be discovered.